

Scalability

Friday, June 4, 2021 8:11 PM

<https://docs.microsoft.com/en-us/azure/architecture/guide/technology-choices/compute-decision-tree#scalability>

Criteria	Virtual Machines	App Service	Service Fabric	Azure Functions	Azure Kubernetes Service	Container Instances	Azure Batch
Autoscaling	Virtual machine scale sets	Built-in service	Virtual machine scale sets	Built-in service	Pod auto-scaling ¹ , cluster auto-scaling ²	Not supported	N/A
Load balancer	Azure Load Balancer	Integrated	Azure Load Balancer	Integrated	Azure Load Balancer or Application Gateway	No built-in support	Azure Load Balancer
Scale limit ³	Platform image: 1000 nodes per scale set, Custom image: 600 nodes per scale set	30 instances, 100 with App Service Environment	100 nodes per scale set	200 instances per Function app	100 nodes per cluster (default limit)	20 container groups per subscription (default limit).	20 core limit (default limit).

Notes

1. See [Autoscale pods](#).
2. See [Automatically scale a cluster to meet application demands on Azure Kubernetes Service \(AKS\)](#).
3. See [Azure subscription and service limits, quotas, and constraints](#).

Availability

Criteria	Virtual Machines	App Service	Service Fabric	Azure Functions	Azure Kubernetes Service	Container Instances	Azure Batch
SLA	SLA for Virtual Machines	SLA for App Service	SLA for Service Fabric	SLA for Functions	SLA for AKS	SLA for Container Instances	SLA for Azure Batch
Multi region failover	Traffic manager	Traffic manager	Traffic manager, Multi-Region Cluster	Azure Front Door	Traffic manager	Not supported	Not Supported

AVAILABILITY

From <<https://docs.microsoft.com/en-us/azure/architecture/guide/technology-choices/compute-decision-tree#scalability>>